
AUTOMATIC SCORING OF COGNITION DRAWINGS

Assessing the quality of machine-based scores against a gold standard

DISCUSSION PAPER

Arne Bethmann
 SHARE Germany
 MEA-SHARE and
 SHARE Berlin Institute
 abethmann@share-berlin.eu

Marina Aoki
 SHARE Germany
 MEA-SHARE and
 SHARE Berlin Institute
 maoki@share-berlin.eu

Charlotte Hunsicker
 SHARE Germany
 MEA-SHARE and
 SHARE Berlin Institute
 chunsicker@share-berlin.eu

Claudia Weileder
 SHARE Germany
 MEA-SHARE and
 SHARE Berlin Institute
 cweileder@share-berlin.eu

Abstract

Figure drawing is often used as part of dementia screening protocols. The Survey of Health Aging and Retirement in Europe (SHARE) has adopted three drawing tests from Addenbrooke’s Cognitive Examination III as part of its questionnaire module on cognition. While the drawings are usually scored by trained clinicians, SHARE uses the face-to-face interviewers who conduct the interviews to score the drawings during fieldwork. This may pose a risk to data quality, as interviewers may be less consistent in their scoring and more likely to make errors due to their lack of clinical training. This paper therefore reports a first proof of concept and evaluates the feasibility of automating scoring using deep learning. We train several different convolutional neural network (CNN) models using about 2,000 drawings from the 8th wave of the SHARE panel in Germany and the corresponding interviewer scores, as well as self-developed ‘gold standard’ scores. The results suggest that this approach is indeed feasible. Compared to training on interviewer scores, models trained on the gold standard data improve prediction accuracy by about 10 percentage points. The best performing model, ConvNeXt Base, achieves an accuracy of about 85%, which is 5 percentage points higher than the accuracy of the interviewers. While this is a promising result, the models still struggle to score partially correct drawings, which are also problematic for interviewers. This suggests that more and better training data is needed to achieve production-level prediction accuracy. We therefore discuss possible next steps to improve the quality and quantity of training examples.

Keywords Automated Scoring · Deep Learning · Convolutional Neural Networks (CNN) · Dementia Screening · Cognitive Drawing Tests · Addenbrooke’s Cognitive Examination III (ACE-III) · Survey Data · Data Quality · Survey of Health Aging and Retirement in Europe (SHARE)

1 Introduction

Dementia and related disorders are among the leading causes of death, disability and dependency among the global elderly population (World Health Organization, 2017). In addition to the severe cognitive, mental and physical impact on individuals with dementia, the disease also affects their social networks and, with increasing prevalence, society as a whole. The estimated number of cases is expected to increase from around 55 million in 2021 to 78 million in 2030 and 139 million in 2050 (World Health Organization, 2021), largely due to the ageing of the world’s population. There are currently no known cures for the disease, making the identification and management of risk factors one of the few promising approaches to address this emerging crisis.

Cognitive tests such as the Montreal Cognitive Assessment (MoCA, see Nasreddine et al., 2005) or the Mini-Mental State Examination (MMSE, see Folstein et al., 1975) can be used to detect early signs of cognitive decline. The correct drawing of figures such as cubes or clocks is often used as an indicator in these tests to assess the visuospatial domain of cognitive functioning. They are usually administered in a clinical setting and used in conjunction with additional assessments to diagnose cognitive impairment, such as dementia-type disorders. The drawings are usually scored by trained clinicians.

More recently, cognitive drawing tests have been used in several large-scale survey studies, such as SHARE, the Survey of Health, Ageing and Retirement in Europe¹ in its main survey, and the HCAP network studies (Langa et al., 2019).² While these tests are adapted from clinical cognitive assessments, in this setting they are used to make population-based estimates, such as the prevalence of dementia symptoms, rather than individual diagnoses. This will allow researchers to look for patterns in early cognitive decline and hopefully identify factors that can slow or delay its progression.

Collecting this data on a large scale in an interviewer-administered survey context means, in the case of SHARE, tens of thousands of interviews conducted repeatedly by thousands of interviewers over several waves. This requires a high degree of standardisation to ensure consistent measurement, which is particularly challenging given that survey interviewers are generally not clinically trained to diagnose cognitive problems.

Since wave 8 (2019/2020), SHARE has included three drawing tests from Addenbrooke’s Cognitive Examination III (see Wagner & Douhou, 2021) as part of the cognitive impairment questionnaire module. To date, SHARE has relied on regular face-to-face interviewers to carry out scoring during fieldwork. Training procedures and office-based scoring have shown that this is not a trivial task and that it warrants closer scrutiny and, ideally, a more standardised approach to mitigate data quality issues.

In this paper, we explore the use of machine learning, specifically convolutional neural networks, to automate cognitive drawing scoring, with the aim of improving data quality and reducing the burden of scoring on interviewers. Our approach involves training models such as ConvNeXt (Liu et al., 2022), AlexNet (Krizhevsky et al., 2012), VGG (Simonyan & Zisserman, 2014) and ResNet50 (He et al., 2015) on two types of datasets: one scored by field interviewers and another using a self-developed ‘gold standard’ procedure. We first assess the ability of the models to learn from interviewer-provided scores, and then examine the improvement in accuracy when these models are trained on the more consistently scored gold standard data. We then compare the accuracy of manual interviewer scoring with model-based scores trained on either dataset. Subsequently, we look at differences in model performance due to different hyperparameter combinations. Finally, we consider the extent and nature of errors produced by the models and suggest how these might be mitigated in further research. This comprehensive analysis aims to demonstrate the potential of deep learning to standardise and improve the accuracy of cognitive assessment scoring in large-scale survey studies.

¹<https://share-eric.eu/>

²<https://hcap.isr.umich.edu/>




Figure 1: Examples of the three drawings: two overlapping infinity loops, a three-dimensional cube, and a clock with its hands set to a specific time.

2 Related Research

Interviewer effects have long been recognised as a potential source of error in survey research and are therefore well studied (see West & Blom, 2016). There is no comprehensive research on interviewer effects in cognitive tests in surveys, but there is reason to believe that interviewers may also influence the data collection process in these areas (Malhotra et al., 2015). Despite additional training, interviewers may face unique challenges when tasked with cognitive tests, as these assessments are not part of their routine responsibilities.

Initial applications of machine learning to cognitive testing have been conducted in clinical neuropsychology to aid in the diagnosis of cognitive impairment (Binaco et al., 2020; Youn et al., 2021). The datasets used in these studies are not only comparatively small, but also very specific in their composition, often including a large proportion of individuals suspected of having some degree of cognitive impairment.

There have also been efforts to use machine learning to classify drawings from cognitive assessments collected in representative surveys. Hu et al. (2022) used data from the National Health and Aging Trends Study to train a machine learning model to score drawings from the Clock Drawing Test (CDT). Similarly, Amini et al. (2021) trained a machine learning model to predict a patient’s dementia status based on drawings from the CDT administered in the Framingham Heart Study.

3 Data & Methodology

In this paper we use data from the eighth wave of the Survey of Health, Ageing and Retirement in Europe (SHARE, Börsch-Supan, 2022).³ SHARE is a biennial panel survey conducted in 28 European countries: all EU Member States except Ireland, plus Israel and Switzerland. The target population is people aged 50 and over, who are asked a wide range of questions about their socio-economic situation, social networks and health. To date, some 600,000 interviews have been collected from around 150,000 respondents in nine waves of the survey. The data are available to the scientific community free of charge.⁴

The actual drawings for the cognitive tests⁵ are made in a recording booklet, which is usually discarded after fieldwork. For this study, we collected the recording booklets from the German SHARE sub-study. Out of a total of 2,878 respondents, 2,374 were asked to complete the drawing tests in the cognitive module of the questionnaire because they were over 60 at the time of the interview. We retrieved and scanned 2,109 booklets.

Cognitive functioning in SHARE

Since its inception, SHARE has had a strong focus on health topics, both physical and mental. For example, items on cognitive functioning have been included since wave 1 (Börsch-Supan & Jürges, 2005,

³The interviewer scores and other data from the survey interviews are part of the SHARE data release, while the drawings used in this paper are not yet published.

⁴<https://share-eric.eu/data/data-access>

⁵see questions CF830 to CF839 in the SHARE generic CAPI questionnaire wave 8: https://share-eric.eu/fileadmin/user_upload/Questionnaires/Q-Wave_8/paperverstion_en_GB_8_2_5b.pdf

Gold standard	Correct	91% (989)	7% (80)	2% (21)
	Partially correct	27% (94)	50% (177)	23% (81)
	Incorrect	2% (7)	23% (74)	75% (245)
		Correct	Partially correct	Incorrect
		Interviewer		

Figure 2: Confusion matrix of interviewer scores against gold standard scores. The percentages in the rows show the proportion of interviewer scores in each class, with correct scores highlighted. Absolute numbers are given in parentheses.

pp. 182–186) and have been collected continuously since then. In wave 8, new cognitive measures were introduced in SHARE, including for the first time tests from Addenbrooke’s Cognitive Examination III (ACE-III, see Hsieh et al., 2013).⁶ These aim to assess a respondent’s abilities in the visuospatial domain of cognitive functioning by asking them to draw three figures: two overlapping infinity loops, a three-dimensional cube, and a clock with its hands set to a specific time (see examples in Figure 1). The loops and the cube had to be copied from a printed picture directly above the drawing area, while the clock had to be drawn from memory (see Wagner & Douhou, 2021, p. 44). These drawings were scored by the interviewers during the interview and the result was recorded in the CAPI questionnaire. For the analyses in this paper, we will focus on the cube drawings. The other drawings will form part of subsequent research.

Interviewer scoring

As part of the fieldwork preparations for each SHARE wave, interviewers receive two days of comprehensive training. During this programme, part of the time is devoted to instructing interviewers in the administration of the cognitive module of the questionnaire. With regard to the scoring of each drawing, the aim was to facilitate the interviewer’s understanding of the specified scoring criteria and to ensure a clear and consistent scoring process. The training materials included a detailed description of the scoring rules as well as a set of sample drawings with scores. The scoring rules for the cube drawing were as follows

- Fully correct copy: the cube has 12 lines, even if the proportions are not perfect.
- Partially correct copy: the cube has fewer than 12 lines, but the general shape of the cube is maintained.
- Incorrect copy

Respondents were allowed to correct mistakes during the drawing or to try again if they wished, but interviewers were not allowed to give feedback or otherwise help respondents with the drawing. Finally,

⁶More recently, the Harmonised Cognitive Assessment Protocol (HCAP, see <https://share-eric.eu/data/data-set-details/share-hcap>) was introduced, an additional in-depth study of cognitive functioning carried out in some of the SHARE countries.

Table 1: Percentage of respondents with a self-reported diagnosis of dementia.

Drawing score	Interviewer	Gold standard
Correct	2.57 %	2.52 %
Parially correct	1.82 %	2.09 %
Incorrect	5.50 %	5.61 %

the drawing of the cube was to be scored by the interviewer immediately after completion. Of the 2,109 recording booklets collected and processed, 1,769 contained usable cube drawings with interviewer scores.

Gold standard scoring

As already known from the literature (e.g. Plank, 2022; Say & O’Driscoll, 2022) and confirmed during the interviewer training sessions and scoring tests in the office, it is clear that scoring the cube drawings is not as straightforward as it might seem. The scoring rules are not always clear and there are many borderline cases. In addition, the interviewers are not clinically trained and may not be able to identify certain types of error. This is particularly true of the partially correct category, where the cube shape is maintained but some lines are missing. For example, it is not always clear whether a line is missing or simply not visible due to the drawing technique. In addition, the interviewers are familiar with the respondents and may include factors other than just looking at the drawing in their judgement and final score. To reduce this measurement error, we also introduce in-house scoring, which we will use as a benchmark and refer to as ‘gold standard’ or ‘ground truth’ scores (depending on the scientific discipline; for the sake of simplicity, we will use ‘gold standard’ for the remainder of the paper). In this process, the scorers are given only the drawings, with no additional information about the respondents, and are asked to score according to the same set of rules as the interviewers.

The gold standard scoring was carried out by three scorers, two of whom had previous experience in scoring drawings from cognitive tests, albeit in the context of the HCAP study, which has a different set of rules. The third scorer had no previous experience. All three were trained in-house using the same materials as the interviewers. Each cube was scored independently by all three scorers. The scorers were instructed to score the drawings according to the same rules as the interviewers, but without access to the previous scores. In cases of disagreement between the scores, including those of the interviewers, an arbitration round was added in which each case was discussed and a final scoring decision was made, by majority vote if necessary. This was done by a team of two of the scorers and a third person. In total, 1,776 cubes were scored by the scorers, slightly more than by the interviewers, as some usable drawings with missing interviewer scores could be recovered. 905 of these had to be decided by the arbitration team. The final gold standard scores were then used to train the models.

A comparison of the interviewer scores with the gold standard shows an accuracy rate of 79.8% (Figure 2). The class percentages in the gold standard are: ‘correct’ 61.65%, ‘partially correct’ 19.91% and ‘incorrect’ 18.44%, so quite unbalanced with about 60/20/20. Scoring correct cubes was the easiest for the interviewers, leading to 91% correct scores in this class. Identifying incorrect cubes was more difficult, but interviewers still achieved an accuracy rate of 75%. The partially correct cubes are indeed problematic, and interviewers achieved only 50% accuracy, still better than chance in our particular 3-class problem, but not very good either.

Table 1 relates the score to the respondent’s self-reported dementia diagnosis as a plausibility check. Although the number of respondents with self-reported dementia is very small, there is some association with incorrectly drawn cubes as opposed to correct cubes, which is what we would hope to see with an indicator in a dementia screening test. Partially correct cubes are actually less likely to be associated with reported dementia than fully correct cubes. This seems a bit strange, but given the small sample size it may just be due to low statistical power. We also see that there does not seem to be a substantial difference between these patterns of association for the interviewer scores and the gold standard scores.




Figure 3: Performance of different model architectures on interviewer scored labels (semi-transparent) and gold standard scored labels (solid). Mean accuracy values across all experiments.

Models

We trained several different computer vision deep learning models to predict the class of cubes drawn by respondents as either ‘correct’, ‘partially correct’ or ‘incorrect’. For the actual training datasets, the drawings were scanned, cropped, resized to 128 by 128 pixels and converted to grey scale. Where this could not be done automatically, manual corrections were made. The training/validation dataset split was 75%/25%, i.e. the actual training dataset size was 1,332 out of a total of 1,776 cases with gold standard labels.

As a general baseline, we used AlexNet (Krizhevsky et al., 2012) and then tested ResNet50 (He et al., 2015) and VGG19_BN (Simonyan & Zisserman, 2014) as intermediate, commonly used models. Finally, we also trained more recent ConvNeXt models (Small, Base and Large, see Liu et al., 2022), which have shown competitive performance in image classification tasks⁷ and which we expected to perform well for our classification problem. All models used were Convolutional Neural Networks (CNN), although we also considered using Vision Transformers (ViT, Dosovitskiy et al., 2020). However, as ViTs are considered to be most useful on larger datasets, we decided to stick with CNNs for the time being. We will revisit this decision in future research.

Experiments

For each model, we ran a series of fine-tuning experiments on both training datasets, the one labelled with the interviewer scores and the one labelled with the gold standard scores. We used models pre-trained on the ImageNet data (Deng et al., 2009). The hyperparameters we varied in the experiments were training duration (50 vs. 100 epochs) and data augmentation (none, transform only, transform plus resize). Each experiment was run three times with different seeds, which were the same for all combinations of hyperparameters. In total, we ran 216 experiments, which became part of our analysis. Model training

⁷see e.g. <https://github.com/huggingface/pytorch-image-models/blob/main/results/README.md>

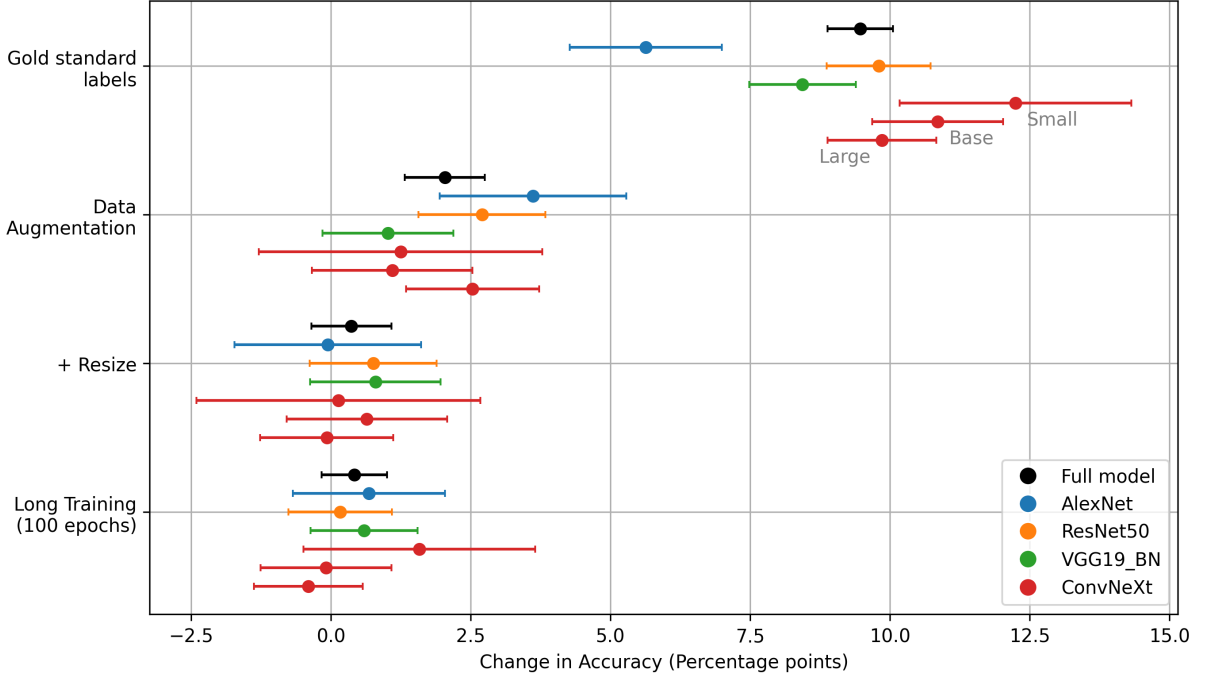


Figure 4: Effect of hyperparameter variation on accuracy across experiments. Dots represent point estimates with error bars showing 95 percent confidence intervals. Reference categories are: interviewer-scored data, no augmentation, and 50 epochs of training.

was performed using the fastai (Howard & Gugger, 2020) and timm (Wightman, 2019) libraries. All experiments were run on an NVIDIA RTX A5000 in a Lenovo P360 workstation with 64 GB of RAM.

4 Results

Figure 3, clearly shows that using our gold standard procedure to label the training data substantially increases prediction accuracy across all model architectures (compare solid circles for gold standard scores vs. semi-transparent circles for interviewer scores). This is likely due to more consistent scoring, i.e. similar looking cubes are less ambiguously scored as the same class. Whether this is actually an indicator of higher quality labels needs to be evaluated by experts such as clinical psychologists. For now, we accept this assumption and see that the increase in accuracy is in the order of ten percentage points for all models except AlexNet. It is also the only model that does not beat the accuracy of human interviewers, which is close to 80%.

An interesting observation is that the difference between the models is more pronounced with higher quality training data. While ConvNeXt performs similarly to ResNet50 and VGG19_BN on the lower quality data, it outperforms all other models when trained on the gold standard data. The amount of computation required is clearly, but not perfectly, related to performance. In our experiments, ConvNeXt Base performed best, even though it is both smaller and less computationally intensive than VGG19_BN and ConvNeXt Large. Even the smaller and less computationally intensive models are competitive. ConvNeXt Small comes second in terms of accuracy, and even ResNet50 is only beaten by the much newer ConvNeXt family of models.

Although this is a rather arbitrary comparison, it is clear that the ConvNeXt models perform well and substantially better than human interviewers, even on a surprisingly small dataset. In future research we will use much larger datasets, which may require a re-evaluation, but for the time being we will use ConvNeXt as our model architecture of choice for further analysis.

True	Correct	98%	2%	0%
	Partially correct	35%	54%	11%
	Incorrect	4%	13%	83%
		Correct	Partially correct	Incorrect
		Predicted		

Figure 5: Confusion matrix of the best ConvNeXt Base model trained with data augmentation for 100 epochs on the gold standard data. Percentages in rows show the proportion of model predictions in each class with correct predictions highlighted. Parentheses contain absolute numbers.

Impact of hyperparameters

To analyse the impact of the hyperparameters, we estimated linear regression models for all CNN model architectures individually, with prediction accuracy as the dependent variable and the hyperparameters as the independent variables. In addition, we estimated a model on the full dataset of experimental results where we also controlled for the CNN model architecture to get a sense of the architecture-independent effect of the hyperparameters.

Figure 4 shows that doubling the training time from 50 to 100 epochs had no significant effect on training accuracy, which may be due to the models quickly adapting to such a small dataset size. Data augmentation had a relatively modest but still significant effect of 2 percentage points in the full model. This effect varied considerably between model architectures, with AlexNet benefiting the most. The addition of the resize variant did not significantly improve prediction accuracy. As seen in the previous figure, the effect of the gold standard is significant at around 9.5 percentage points. AlexNet only gets a comparatively small boost of 5.5 percentage points, while ConvNeXt improves by about 12 percentage points compared to when trained on the interviewer-scored labels. However, both have a relatively wide confidence interval.

Classification of errors

As mentioned above, not all cubes are equally difficult to score. In particular, the partially correct cubes are problematic as they are often incorrectly scored by interviewers (see Figure 2). We therefore looked more closely at the errors made by the models trained on the gold standard data.

Figure 5 shows the confusion matrices for the ConvNeXt Base model. The model performs best on the correct cubes, with an accuracy of 98%. Incorrectly drawn cubes are classified as such with an accuracy of 81%. As we saw with the interviewers, the partially correct cubes are the most difficult to classify, with an accuracy of 73%. This is certainly better than the interviewer performance, but still much worse than the other classes.

The model we used for comparison outperforms the interviewers in all classes by a considerable margin, but still struggles with the partially correct cubes. One reason for this may be that, as an intermediate category, it is inherently ambiguous and thus harder to classify than the more clear-cut correct and




Figure 6: Production plot of cumulative accuracy over prediction confidence per class. Predictions are based on the best ConvNeXt Base model, trained with data augmentation for 100 epochs on the gold standard data.

incorrect categories. This also makes it difficult to assign correct labels during the gold standard scoring process. The fact that the scoring rules are so vague exacerbates the problem, which is why it seems crucial to rely on expert judgement in the future. The small sample size of partially correct cubes in the gold standard data may also be a factor. This is probably also true for the incorrect class, where accuracy is higher but not outstanding.

Automated scoring in production

If we are going to use automated scoring based on model predictions to improve the quality of the data in the SHARE data releases made available to substantive researchers, we need to know what proportion of the cubes can be automatically labelled while still maintaining a level of accuracy that we consider acceptable. To get an idea of this relationship, we ordered the cubes by the confidence of the prediction, i.e. the probability of the most likely class, on the x-axis, with the most confident on the left and the least confident on the right. We then plotted this against the cumulative accuracy of the model predictions on the y-axis.⁸ To distinguish between different types of error, we also plotted the cumulative accuracy of the correct, partially correct and incorrect cubes separately. The result is shown in Figure 6.

What we see in the plot is that for predictions with very high confidence, the accuracy is also very high, but that the accuracy drops off rapidly as the confidence decreases along the x-axis. On the far right we end up with the same percentages we saw in Figure 5, which makes sense as this would be the case if we just let the model automatically score all the cubes. If we want higher accuracy we need to move to the left, but this comes at the cost of having to score more cubes manually. For example, if we wanted to have an overall accuracy of about 95%, we would need to manually score about 20% of the cubes and let the model score 80%, which is still a significant reduction in the scoring burden. If we are willing to accept an overall accuracy of 90%, we can reduce the scoring burden to about 5%. Depending on how we value the errors in the different classes, we can place more emphasis on the class-specific curves. For

⁸see also Bethmann et al. (2014), where we used a similar approach for automated coding of occupations in survey data

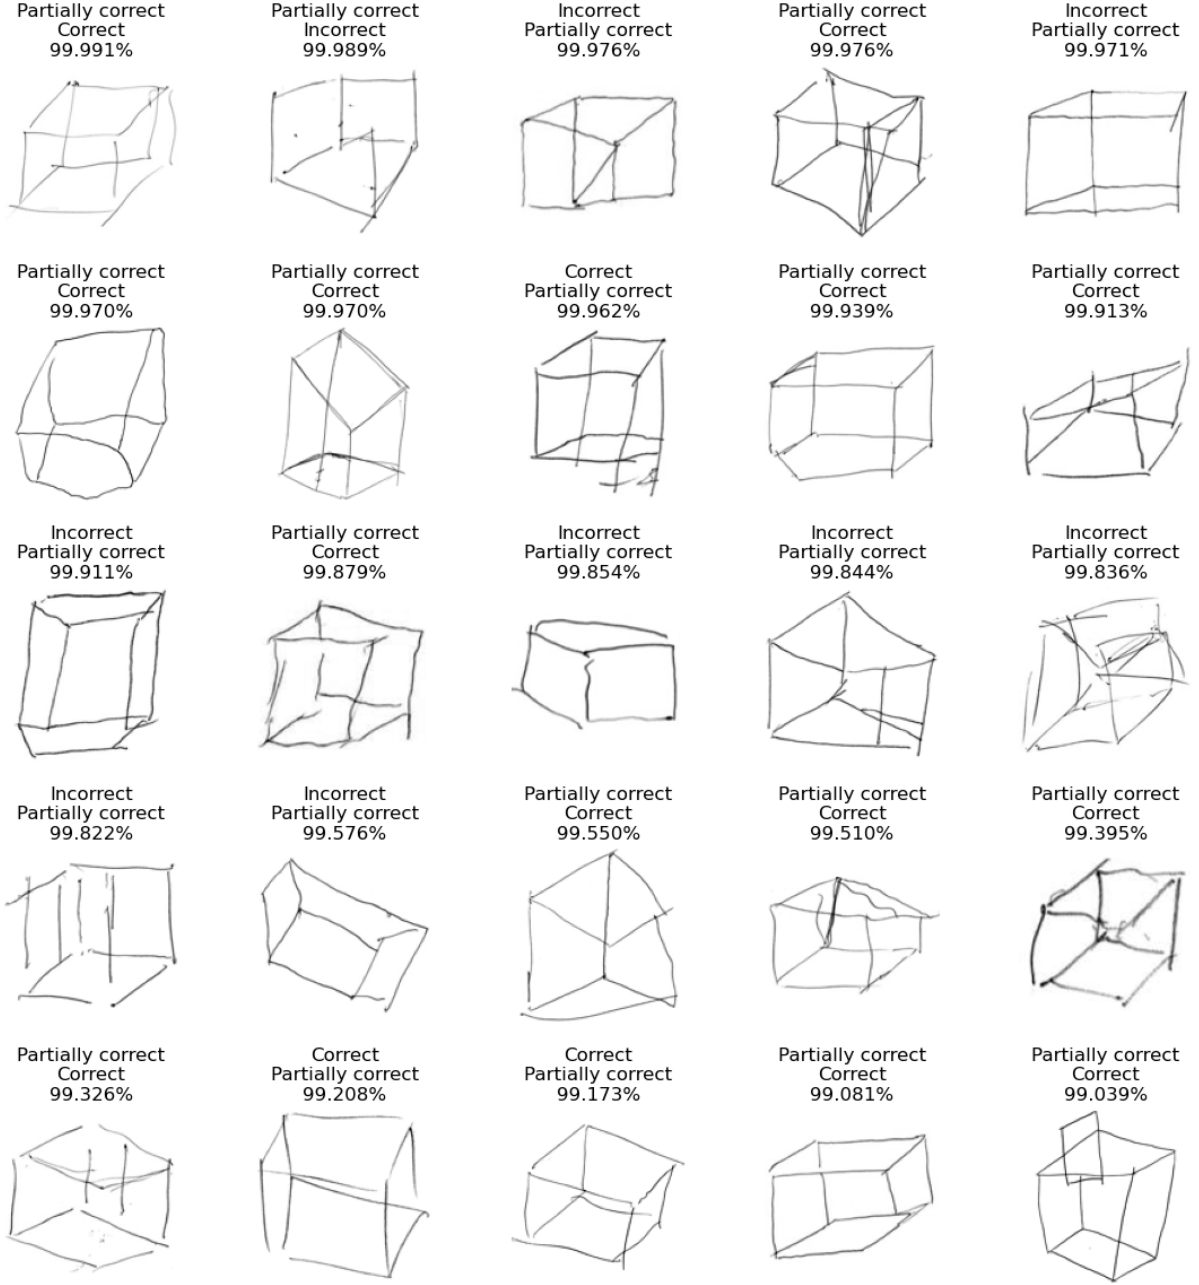


Figure 7: Misclassified cubes with highest confidence. True class, predicted class and prediction confidence are given above each drawing.

example, if we wanted to have an accuracy of about 85% for the partially correct cubes, we would need to score about 40% of the cubes manually and let the model score 60%.

There are some caveats to these graphs. They only give an estimate of the expected accuracy, as confidence and accuracy are based on the model's prediction, and as we can see, the model itself is far from perfect. This is indicated not only by the accuracies being (well) below 100%, but also by the jaggedness of the downward slope, which could be due to errors in labelling the data, or the model struggling to classify certain types of cubes. In either case, the model made an incorrect classification even though it was very confident about the prediction (see Figure 7 for a few examples). This is obviously a problem if the model is to be used in production, and should be mitigated by improving both the training data and the model (see some suggestions in the next section).

5 Discussion

In conclusion, we have shown that it is indeed possible to use deep learning to automate the scoring of drawing tests in cognitive assessments collected in surveys. In this proof of concept, we trained several different CNN models on two types of datasets: one labelled with scores provided by survey interviewers during fieldwork, and another labelled with an in-house ‘gold standard’ scoring procedure. The models in the ConvNeXt family averaged over 85% overall prediction accuracy in our experiments when trained on the gold standard data, meaning they outperformed human interviewers by around 5 percentage points. This is a promising result, especially considering that it can be achieved without any further human intervention, thus reducing the burden on the interviewer while improving the quality of the scoring at virtually no additional cost.

Limitations

We would be reluctant to use the current models in production at this stage. The main reason for this is that the models, like the interviewers, are still struggling to score the partially correct cubes. They have not yet achieved an acceptable level of accuracy in this class. This is probably due to the small number of cases in the training dataset (about 350) and possibly also to ambiguities in the gold standard scores. At the same time, we were quite surprised to achieve this level of performance with a training dataset of only about 1,300 cases in total after subtracting the validation dataset.

Increasing the quantity of training data

In terms of the volume of cube drawings, we have already collected around 55,000 additional recording booklets from SHARE waves eight and nine in several countries. Once these have been scanned and processed, we will have a much larger dataset on which to train our models. The caveat here is that so far we only have interviewer scores for these cubes, which, as we have shown, are not a particularly good basis for training. We therefore need to score these cubes using our gold standard procedure.

As this is a very time-consuming process, we will need to adapt the procedure to make it more efficient. To do this, we will look at automatic pre-labelling of the cubes using the models trained on the current dataset, as well as active learning approaches such as uncertainty sampling (e.g. Cohn et al., 1994; Settles, 2012). This will allow us to focus on the more difficult, and therefore useful, cases and thus speed up the process.

While the larger dataset is likely to mitigate some of the detrimental effects of class imbalance on model performance, in particular the small number of partially correct cubes, we will also want to explore the possibility of using data augmentation to generate more of these cases. A particularly interesting approach could be to use diffusion models to generate synthetic data, which have recently been shown to be very effective in improving performance in image classification tasks (e.g. Azizi et al., 2023; Schwag, 2023).

Improving the quality of the scoring

We will also need to rethink the assumption that our ‘gold standard’ scores are not only more consistent, and therefore easier to train on, but actually more accurate than interviewer scores. This means that for at least a representative subset of the gold standard scores we will need to obtain expert judgements from, for example, clinical psychologists, something that we might actually want to incorporate into our use of active learning approaches.

It will also be interesting to look more closely at the scoring rules and the scoring process itself. As we have seen, the partially correct cubes are particularly problematic and the scoring rules are not very clear. It seems questionable whether the scoring rules in their current form are actually a good way of communicating the scoring criteria to the scorers. At the moment there seems to be a lot of hidden knowledge that scorers have to acquire through experience, which may work well if the scorers are true

experts, such as clinical psychologists, who can apply their knowledge gained through clinical training to the scoring process. This does not seem feasible for ordinary survey interviewers or lay office scorers, who would have to rely much more on scoring rules, training examples and expert supervision. We will therefore look more closely at what different types of scorers actually do, what their inter-rater reliabilities are, and how this relates to quality criteria such as expert ratings, other test scores and diagnoses. One way to inform these analyses is to use the trained CNN models to visualise the parts of the drawings that are most important for classification with techniques such as Grad-CAM (Selvaraju et al., 2016). This will give us a better idea of how the scoring rules currently work and how they could be improved, at least if we want to continue using lay scorers to score cognitive drawing tests in the near future.

Respondents draw cubes in many different ways and the deviations from the template pictures are very varied. These deviations are often not random, but follow certain patterns that appear in many drawings, for reasons that may include situational factors, personal style, practice, cultural differences, physical impairments or, of course, different levels of cognitive functioning in the visuospatial domain. At the same time, it is often not clear, especially to a lay scorer, whether a particular deviation should affect the score or not. The scoring rules are far too short and vague to capture this level of detail. It can therefore be useful to use dimensionality reduction or clustering techniques to get an idea of the types of deviation patterns present in the distribution (Ester et al., 1996; Maaten & Hinton, 2008; McInnes et al., 2018). The combination of expert judgement and model-based accuracy predictions could then be used to identify relevant clusters, for example because they are ambiguously scored or because they are a particularly salient indicator of cognitive impairment. This could help to guide the labelling process beyond the creation of more complex but generic rules (as is done in other cognitive assessments) and expert supervision of scoring.

Improving the models

There is also plenty of room for improvement on the modelling side, including but not limited to more hyperparameter tuning.

The ImageNet pre-training weights we used, as provided by the fastai/timm default option, may not be the best, or even the most obvious, basis for our fine-tuning. The dataset consists mainly of photographs, which are probably not ideal for pre-training the layers of a CNN tasked with classifying human-drawn sketches of very simple objects. We will therefore explore more suitable pre-training datasets in future research, such as ImageNet-Sketch (Wang et al., 2019), TU Berlin Sketch Dataset (Eitz et al., 2012), QuickDraw (Ha & Eck, 2017), or perhaps just the cumulated set of cube, infinity loop, and clock drawings from our own data.

In terms of model architectures, we will also explore the use of Vision Transformers (Dosovitskiy et al., 2020), which are currently considered state-of-the-art for many image classification tasks. We expect them to be a viable option, especially as dataset size increases, potentially challenging the performance of the best CNNs.

Future work

As a future step, we will score the other two drawings from the ACE-III test, the overlapping infinity loops and the clock drawing, using the best practices learned from scoring the cube drawings. We will also explore the possibility of using the models trained on the cube drawings to pre-label the other drawings, which should speed up the process considerably.

An interesting avenue to explore is transfer learning from the cube drawings to other drawings. This could be done by simply fine-tuning the cube-trained models with the training datasets of, say, the infinity loops or the clocks. The rationale would be that the features learned during training on the cube dataset are more easily tuned to the classification task of other similar drawings than the original pre-training weights based on, for example, ImageNet. Similarly, we could explore transfer learning from our models

to cognitive tests in other studies, such as those used in HCAP, which include very similar drawing tests. This could be particularly interesting since, for example, SHARE-HCAP is a sub-sample of about 2,500 respondents from the SHARE sample, and fine-tuning a good model for a similar task might make it possible to achieve production-level prediction accuracy even with such small training datasets.

Finally, we also look forward to working on making our data and models available to the scientific community. We plan to make the drawings and the corresponding scores available as open datasets. We also want to release the models and the code used to train them, so that other researchers can replicate our results and build on them in their own research. Researchers will also be able to use the models to score their own drawings, which we hope will lead to more standardised and consistent scoring of cognitive drawing tests in future research.

Acknowledgments

This paper uses data from SHARE Wave 8 (DOI: 10.6103/SHARE.w8.800) see Börsch-Supan et al. (2013) for methodological details. The SHARE data collection has been funded by the European Commission, DG RTD through FP5 (QLK6-CT-2001-00360), FP6 (SHARE-I3: RII-CT-2006-062193, COMPARE: CIT5-CT-2005-028857, SHARELIFE: CIT4-CT-2006-028812), FP7 (SHARE-PREP: GA N°211909, SHARE-LEAP: GA N°227822, SHARE M4: GA N°261982, DASISH: GA N°283646) and Horizon 2020 (SHARE-DEV3: GA N°676536, SHARE-COHESION: GA N°870628, SERISS: GA N°654221, SSHOC: GA N°823782, SHARE-COVID19: GA N°101015924) and by DG Employment, Social Affairs & Inclusion through VS 2015/0195, VS 2016/0135, VS 2018/0285, VS 2019/0332, and VS 2020/0313. Additional funding from the German Ministry of Education and Research, the Max Planck Society for the Advancement of Science, the U.S. National Institute on Aging (U01_AG09740-13S2, P01_AG005842, P01_AG08291, P30_AG12815, R21_AG025169, Y1-AG-4553-01, IAG_BSR06-11, OGHA_04-064, HHSN271201300071C, RAG052527A) and from various national funding sources is gratefully acknowledged (see www.share-project.org).

References

- Amini, S., Zhang, L., Hao, B., Gupta, A., Song, M., Karjadi, C., Lin, H., Kolachalama, V. B., Au, R., & Paschalidis, I. C. (2021). An artificial intelligence-assisted method for dementia detection using images from the clock drawing test. *Journal of Alzheimer's Disease : JAD*, 83(7, 2), 581–589. <https://doi.org/10.1017/S1355617720000144>
- Azizi, S., Kornblith, S., Saharia, C., Norouzi, M., & Fleet, D. J. (2023). *Synthetic data from diffusion models improves ImageNet classification*. <https://doi.org/10.48550/ARXIV.2304.08466>
- Bethmann, A., Schierholz, M., Wenzig, K., & Zielonka, M. (2014). Automatic coding of occupations. *Proceedings of Statistics Canada Symposium 2014. Beyond Traditional Survey Taking: Adapting to a Changing World*. <https://www.statcan.gc.ca/sites/default/files/media/14291-eng.pdf>
- Binaco, R., Calzaretto, N., Epifano, J., McGuire, S., Umer, M., Emrani, S., Wasserman, V., Libon, D. J., & Polikar, R. (2020). Machine learning analysis of digital clock drawing test performance for differential classification of mild cognitive impairment subtypes versus alzheimer's disease. *Journal of the International Neuropsychological Society*, 26(7), 690–700. <https://doi.org/https://doi.org/10.1017/S1355617720000144>
- Börsch-Supan, A. (2022). *Survey of health, ageing and retirement in europe (SHARE) wave 8*. SHARE-ERIC. <https://doi.org/10.6103/SHARE.W8.800>
- Börsch-Supan, A., Brandt, M., Hunkler, C., Kneip, T., Korbmacher, J., Malter, F., Schaan, B., Stuck, S., & Zuber, S. (2013). Data resource profile: The survey of health, ageing and retirement in europe (SHARE). *International Journal of Epidemiology*, 42(4), 992–1001. <https://doi.org/10.1093/ije/dyt088>
- Börsch-Supan, A., & Jürges, H. (Eds.). (2005). *The survey of health, ageing and retirement in europe – methodology*. MEA. https://share-eric.eu/fileadmin/user_upload/Methodology_Volumes/Methodology_2005.pdf
- Cohn, D., Atlas, L., & Ladner, R. (1994). Improving generalization with active learning. *Machine*

- Learning*, 15(2), 201–221. <https://doi.org/10.1007/bf00993277>
- Deng, J., Dong, W., Socher, R., Li, L.-J., Li, K., & Fei-Fei, L. (2009, June). ImageNet: A large-scale hierarchical image database. *2009 IEEE Conference on Computer Vision and Pattern Recognition*. <https://doi.org/10.1109/cvpr.2009.5206848>
- Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., Dehghani, M., Minderer, M., Heigold, G., Gelly, S., Uszkoreit, J., & Houlsby, N. (2020). *An image is worth 16x16 words: Transformers for image recognition at scale*. <https://doi.org/10.48550/ARXIV.2010.11929>
- Eitz, M., Hays, J., & Alexa, M. (2012). How do humans sketch objects? *ACM Trans. Graph. (Proc. SIGGRAPH)*, 31(4), 44:1–44:10. https://cybertron.cg.tu-berlin.de/eitz/pdf/2012_siggraph_classifysketch.pdf
- Ester, M., Kriegel, H.-P., Sander, J., & Xu, X. (1996). A density-based algorithm for discovering clusters in large spatial databases with noise. *Kdd*, 96(34), 226–231.
- Folstein, M. F., Folstein, S. E., & McHugh, P. R. (1975). "Mini-mental state". *Journal of Psychiatric Research*, 12(3), 189–198. [https://doi.org/10.1016/0022-3956\(75\)90026-6](https://doi.org/10.1016/0022-3956(75)90026-6)
- Ha, D., & Eck, D. (2017). *A neural representation of sketch drawings*. <https://doi.org/10.48550/ARXIV.1704.03477>
- He, K., Zhang, X., Ren, S., & Sun, J. (2015). *Deep residual learning for image recognition*. <https://doi.org/10.48550/ARXIV.1512.03385>
- Howard, J., & Gugger, S. (2020). Fastai: A layered API for deep learning. *Information*, 11(2). <https://doi.org/10.3390/info11020108>
- Hsieh, S., Schubert, S., Hoon, C., Miosh, E., & Hodges, J. (2013). Validation of the addenbrooke’s cognitive examination III in frontotemporal dementia and alzheimer’s disease. *Dementia and Geriatric Cognitive Disorders*, 36(3–4), 242–250. <https://doi.org/10.1159/000351671>
- Hu, M., Murphey, Y., Wang, S., Qin, T., Zhao, Z., Gonzalez, R., Freedman, V., & Zahodne, L. (2022). Exploring the use of machine learning to improve dementia detection: Automating coding of the clock-drawing test. *Innovation in Aging*, 6, 382–382. <https://doi.org/10.1093/geroni/igac059.1508>
- Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). ImageNet classification with deep convolutional neural networks. In F. Pereira, C. J. Burges, L. Bottou, & K. Q. Weinberger (Eds.), *Advances in neural information processing systems* (Vol. 25). Curran Associates, Inc. https://proceedings.neurips.cc/paper_files/paper/2012/file/c399862d3b9d6b76c8436e924a68c45b-Paper.pdf
- Langa, K. M., Ryan, L. H., McCammon, R. J., Jones, R. N., Manly, J. J., Levine, D. A., Sonnega, A., Farron, M., & Weir, D. R. (2019). The health and retirement study harmonized cognitive assessment protocol project: Study design and methods. *Neuroepidemiology*, 54(1), 64–74. <https://doi.org/10.1159/000503004>
- Liu, Z., Mao, H., Wu, C.-Y., Feichtenhofer, C., Darrell, T., & Xie, S. (2022). *A ConvNet for the 2020s*. <https://doi.org/10.48550/ARXIV.2201.03545>
- Maaten, L. van der, & Hinton, G. (2008). Visualizing data using t-SNE. *Journal of Machine Learning Research*, 9(86), 2579–2605. <http://jmlr.org/papers/v9/vandermaaten08a.html>
- Malhotra, R., Haaland, B. A., Chei, C.-L., Chan, A., Malhotra, C., & Matchar, D. B. (2015). Presence of and correction for interviewer error on an instrument assessing cognitive function of older adults. *Geriatrics & Gerontology International*, 15(3), 372–380. <https://doi.org/10.1111/ggi.12331>
- McInnes, L., Healy, J., & Melville, J. (2018). *UMAP: Uniform manifold approximation and projection for dimension reduction*. <https://doi.org/10.48550/ARXIV.1802.03426>
- Nasreddine, Z. S., Phillips, N. A., Bédirian, V., Charbonneau, S., Whitehead, V., Collin, I., Cummings, J. L., & Chertkow, H. (2005). The montreal cognitive assessment, MoCA: A brief screening tool for mild cognitive impairment. *Journal of the American Geriatrics Society*, 53(4), 695–699. <https://doi.org/10.1111/j.1532-5415.2005.53221.x>
- Plank, B. (2022). *The ‘problem’ of human label variation: On ground truth in data, modeling and evaluation*. <https://doi.org/10.48550/ARXIV.2211.02570>
- Say, M. J., & O’Driscoll, C. (2022). Inter-rater variability in scoring of addenbrooke’s cognitive examination-third edition (ACE-III) protocols. *Applied Neuropsychology: Adult*, 1–5. <https://doi.org/10.1080/>

23279095.2022.2083964

- Sehwag, V. (2023). Minimal-diffusion: A minimal yet resourceful implementation of diffusion models (along with pretrained models + synthetic images for nine datasets). *GitHub Repository*. <https://github.com/VSehwag/minimal-diffusion>
- Selvaraju, R. R., Cogswell, M., Das, A., Vedantam, R., Parikh, D., & Batra, D. (2016). Grad-CAM: Visual explanations from deep networks via gradient-based localization. *International Journal of Computer Vision*, 128(2), 336–359. <https://doi.org/10.1007/s11263-019-01228-7>
- Settles, B. (2012). Uncertainty sampling. In *Synthesis lectures on artificial intelligence and machine learning* (pp. 11–20). Springer International Publishing. https://doi.org/10.1007/978-3-031-01560-1_2
- Simonyan, K., & Zisserman, A. (2014). *Very deep convolutional networks for large-scale image recognition*. <https://doi.org/10.48550/ARXIV.1409.1556>
- Wagner, M., & Douhou, S. (2021). Cognitive measures. In M. Bergmann & A. Börsch-Supan (Eds.), *SHARE wave 8 methodology: Collecting cross-national survey data in times of COVID-19* (pp. 43–46). MEA, Max Planck Institute for Social Law; Social Policy. https://share-eric.eu/fileadmin/user_upload/Methodology_Volumes/SHARE_Methodenband_WEB_Wave8_MFRB.pdf
- Wang, H., Ge, S., Xing, E. P., & Lipton, Z. C. (2019). Learning robust global representations by penalizing local predictive power. *NeurIPS 2019*. <https://doi.org/10.48550/ARXIV.1905.13549>
- West, B. T., & Blom, A. G. (2016). Explaining interviewer effects: A research synthesis. *Journal of Survey Statistics and Methodology*, 5(2), 175–211. <https://doi.org/10.1093/jssam/smw024>
- Wightman, R. (2019). PyTorch image models. In *GitHub repository*. <https://github.com/rwightman/pytorch-image-models>; GitHub. <https://doi.org/10.5281/zenodo.4414861>
- World Health Organization. (2017). *Global action plan on the public health 2005 - 2007*. <https://www.who.int/publications/i/item/9789241513487>
- World Health Organization. (2021). *Global status report on the public health response to dementia*. <https://www.who.int/publications/i/item/9789240033245>
- Youn, Y. C., Pyun, J.-M., Ryu, N., Baek, M. J., Jang, J.-W., Park, Y. H., Ahn, S.-W., Shin, H.-W., Park, K.-Y., & Kim, S. Y. (2021). Use of the clock drawing test and the rey-osterrieth complex figure test-copy with convolutional neural networks to predict cognitive impairment. *Alzheimer's Research & Therapy*, 13(1), 85. <https://doi.org/10.1186/s13195-021-00821-8>